

LiveTradeBench: Seeking Real-World Alpha with LLMs

Evaluating Large Language Model Agents in Evolving Financial Markets

Haofei Yu, Fenghai Li, Jiaxuan You

University of Illinois, Urbana-Champaign (UIUC)

arXiv:2511.03628 (Nov 2025)

Focus Questions & Key Topics

- Why do static LLM benchmarks (MMLU, GSM8K, Chat LMArena) fail to evaluate **sequential decision-making under uncertainty**?
- What are the **3 core design principles** of the LiveTradeBench environment?
- How do **U.S. Stock Markets** and **Polymarket Prediction Markets** differ as evaluation testbeds?
- What trading performance metrics (Cumulative Return, Sharpe Ratio, Max Drawdown) emerge across **21 LLMs evaluated over 50 live trading days**?
- Why does a high Chat benchmark score **not imply superior trading returns (Alpha)**?
- What behavioral archetypes (**Defensive vs. Aggressive**) and reasoning dynamics do LLMs display?

Motivation: Beyond Static Evaluation

- **Limitations of Existing LLM Benchmarks:**

- Knowledge quizzes and math tests evaluate models on *static, single-turn inputs*.
- Interactive web/computer agents operate in *controllable, deterministic environments*.

- **The Financial Trading Challenge:**


- Financial markets are **continuous, autonomous, and non-stationary**.
- The world evolves independently of the agent; actions adjust portfolio allocations under **real-time live uncertainty**.
- Demands sequential decision-making, risk management, and cross-asset reasoning.

LiveTradeBench vs. Existing Trading Benchmarks

Benchmark	Sequential Decision	Portfolio Management	Live Data Test	Multi-Market
FinQA (Chen et al.)	✗	✗	✗ (Static)	✗ (Single)
FLUE (Shah et al.)	✗	✗	✗ (Static)	✗ (Single)
FinAgentBench (Bigéard et al.)	✗	✗	✗ (Static)	✗ (Single)
StockBench (Chen et al.)	1-Asset	✗	✗ (Backtest)	✗ (Single)
LiveTradeBench (Ours)	✓	✓ (Multi-Asset)	✓ (Live Stream)	✓ (Dual)

LiveTradeBench Design & Execution Pipeline

LiveTradeBench Architecture & Execution Pipeline



Observation & Action Space Formulation

- **Observation Space** $o_t = (q_t, p_t, c_t)$:
 - **Position** q_t : Continuous holdings across all assets + CASH.
 - **Market Price** p_t : Real-time price stream for equities and binary contracts.
 - **Market Context** c_t : Live news feeds collected from Google News.
- **Action Space** a_t (**Portfolio Allocation Abstraction**):
 - Output percentage weights $a_t = [a_t^{(1)}, \dots, a_t^{(K)}, a_t^{(\text{CASH})}]$ where $\sum a_i = 1.0$.
 - Direct mapping from high-level allocation to executable **BUY / SELL / HOLD** trades without low-level execution complexity.

Dual Market Environments

U.S. Stock Market

15 Equities + CASH

- Tech: AAPL, MSFT, NVDA, META
 - Finance & Ind: JPM, V, XOM, CAT
 - Consumer & Health: AMZN, PG, UNH
-

Market Characteristics

Mature, institutional, smooth price trends,
rewards structural risk diversification.

Polymarket Prediction

10 Binary Contracts + CASH

- "Fed rate cut in 2025?"
 - "US Recession in 2025?"
 - "Tether insolvent / USDT depeg?"
-

Market Characteristics

Decentralized, high volatility, sharp jumps,
rewards rapid event-driven belief updating.

Live Test Experimental Setup

- **50-Day Live Test Period:** August 18, 2025 to October 24, 2025.
- **21 Evaluated LLMs Across 6 Major Model Families:**
 - **OpenAI:** GPT-4o, GPT-4.1, GPT-5, GPT-o3
 - **Anthropic:** Claude-3.5-Sonnet, Claude-3.5-Opus, Claude-4, Claude-4.1
 - **xAI:** Grok-3, Grok-4
 - **Alibaba:** Qwen-2.5-72B, Qwen-2.5-235B, Qwen-2.5-235B-Thinking
 - **Meta / DeepSeek / Moonshot:** Llama-3.3, Llama-4, DeepSeek-V3.1, DeepSeek-R1, Kimi K2
- **5 Evaluation Metrics:** Cumulative Return (CR), Sharpe Ratio (SR), Maximum Drawdown (MDD), Win Rate (WR), Volatility (Vol).

Finding 1: Disconnect Between Chat Scores & Alpha

- **High Chat Scores $\neq$ Superior Trading Alpha:**
 - Models with top rankings on **LMarena** (e.g., GPT-4o, Claude-3.5-Sonnet) do *not* automatically generate positive trading returns.
- **Why the Disconnect?**
 - Chat benchmarks reward conversational fluency, whereas trading rewards **discipline, risk control, and dynamic adaptation**.
 - High-scoring chat LLMs often **over-trade**, react to noise, or hallucinate short-term sentiment trends.

Finding 2: LLM Portfolio Archetypes & Cash Dynamics

Defensive Archetype

(DeepSeek-R1, Claude-4)

CASH: 80% - 90%

Assets

- Low Drawdown, Minimal Risk
- Preserves capital during market drops
- May miss rapid upside rallies

Aggressive Archetype

(Grok-3, GPT-5, Qwen-235B)

CASH

High Asset Allocation: 85% - 95%

- High Return Potential & High Volatility
- Highly responsive to live news signals
 - Subject to larger drawdowns

Finding 3: The Cross-Market Generalization Gap

- **Market Non-Generalizability:**

- A model's trading performance on U.S. Stocks does **not** correlate with its performance on Polymarket prediction markets.

- **Key Structural Differences:**

- **U.S. Stocks:** Mature institutional consensus; rewards long-term structural analysis and sector diversification.
- **Polymarket:** High volatility, speculative sentiment, sharp discrete price jumps; rewards **rapid event-driven belief updating**.

Rebalancing Frequency & News Sensitivity

- **Rebalancing Interval (k):**
 - Rebalancing too frequently ($k = 1$ day) leads to **over-reaction to market noise** and transaction penalty.
 - Rebalancing too slowly ($k = 5$ days) fails to adapt to sudden market inflections.
- **News Sensitivity:**
 - Live news signals act as critical catalysts. Models that integrate news context exhibit superior risk-adjusted returns (higher Sharpe Ratio).

Case Study: Belief-Based Reasoning in Action

- **Reasoning Dynamics (e.g., DeepSeek-R1 & Grok-3):**
 - When fed news of unexpected Fed rate shifts, **DeepSeek-R1** explicitly re-evaluates risk metrics in its reasoning trace before altering allocations.
 - **Grok-3** dynamically shifts capital between Tech stocks and CASH based on macro employment data signals.
- **Conclusion:** Explicit Chain-of-Thought (CoT) reasoning improves decision consistency under market uncertainty.

Concept Check Question 1

What is a key design principle of LiveTradeBench that distinguishes it from traditional offline backtesting benchmarks?

- **A.** It uses historical static price datasets from 2010.
- **B.** It streams live market prices and news in real-time, eliminating data leakage and offline backtesting bias.
- **C.** It limits agent actions to single-asset discrete buy/sell commands.
- **D.** It evaluates models exclusively on single-turn conversational quizzes.



Concept Check Question 1: Answer

What is a key design principle of LiveTradeBench that distinguishes it from traditional offline backtesting benchmarks?

- **Correct Answer: B**
- **Explanation:** LiveTradeBench streams live real-time market data and news, preventing offline data leakage and evaluating LLMs under true live uncertainty.



Concept Check Question 2

What did the 50-day live evaluation of 21 LLMs reveal regarding Chat benchmark rankings (LMArena) and real-world trading returns?

- **A.** High LMArena scores perfectly predict high trading Alpha.
- **B.** High LMArena scores do not imply superior trading returns, as chat metrics ignore risk control and sequential decision-making.
- **C.** All LLMs achieved identical trading returns regardless of model family.
- **D.** Stock market performance perfectly generalizes to prediction market performance.



Concept Check Question 2: Answer

What did the 50-day live evaluation of 21 LLMs reveal regarding Chat benchmark rankings (LMArena) and real-world trading returns?

- **Correct Answer: B**
- **Explanation:** Chat benchmarks evaluate single-turn conversational fluency, which fails to capture real-world risk management, asset allocation, and decision consistency.



Recap: Fill-in-the-blank Quiz

Test your understanding of the core findings in LiveTradeBench:

1. LiveTradeBench formulates agent decisions as multi-asset management rather than single-asset buy/sell orders.
2. The benchmark evaluates agents across two distinct environments: U.S. and Polymarket markets.
3. Models like DeepSeek-R1 adopt a portfolio style by holding high CASH percentages to minimize drawdown.
4. Evaluating LLMs in live continuous markets exposes the gap between static benchmark scores and real-world .



Recap: Answers & Summary

Here are the completed concepts:

1. LiveTradeBench formulates agent decisions as multi-asset **portfolio** management rather than single-asset buy/sell orders.
2. The benchmark evaluates agents across two distinct environments: U.S. **stocks** and Polymarket **prediction** markets.
3. Models like DeepSeek-R1 adopt a **defensive (cautious)** portfolio style by holding high CASH percentages to minimize drawdown.
4. Evaluating LLMs in live continuous markets exposes the gap between static benchmark scores and real-world **competence (Alpha)**.



Conclusion & Key Takeaways

- **Live Benchmarking Necessity:**
 - Evaluating LLM agents requires dynamic, non-stationary live environments to test true sequential decision-making under uncertainty.
- **Portfolio Management Abstraction:**
 - Multi-asset allocation (a_t) naturally forces models to balance risk vs. return and manage cash reserves.
- **Future Directions:**
 - Integrating explicit risk constraints, multi-agent communication, and long-horizon memory in financial LLM architectures.